

The Eyes of Iorek Byrnison

John Oberlin
Brown University, 2014

Abstract

State of the art techniques in object detection and pose estimation are powerful and general but usually run at a rate less than 1 Hz. This makes it difficult to employ such techniques in real-time human-computer interaction. This document outlines a simple, robust framework for object detection which trades a large memory overhead for improvements in latency and total throughput of detections. Included is a workflow for that framework which makes training and calibration first intuitive and then automatic.

1 Introduction

Our overall pipeline can be described as:

1. Collect RGB-D data for objects.
2. Train BoW model and kNN classifiers for objects.
3. Use the classifiers and additional logic to provide 3D detections pose estimates of objects.

2 Basic Techniques

Objectness try clustering the proposals for blueboxes.

Fast Keypoints

SIFT descriptors

KMeans

BoW

Color Histogram

Depth Histogram yet to implement

kNN

3 Detector Structure

Green Boxes

Blue Boxes

Red Boxes

4 Teaching the System

This is teaching not training because it is an interaction used to collect data not a script used to optimize a function.

4.1 Background Filtering

Teaching is carried out on a flat surface of uniform color. The background plane and color model are inferred by looking at parts of the image not contained in the blue boxes. These are used to keep with high probability only keypoints which are on the object. During inference (employment), background contributions will contribute to all examples approximately equally, and any confusion is likely to be understandable.

4.2 Manual Teaching

Teaching the system about an object involves showing it many different views of the object in a controlled environment. Once a session is initiated, the following process should be performed:

1. Repeat k times:
 - (a) Adjust the object to expose a novel viewpoint and record pose data.
 - (b) Remove extraneous objects in the scene to ensure a unique blue box is present.
 - (c) Collect the view, which saves RGB-D data and the background data.

Where k might be around 100. Now the collected data can be used to train object class and pose classifiers.

4.3 Self Filtering

Suppose that we want to gather data while a robot is holding the object. It is possible to impose geometric constraints on kept data to ensure that captured keypoints belong to the object being examined and not to our manipulating robot. This is called self filtering.

4.4 Semi-Automatic Teaching

1. Repeat k times:
 - (a) The object is placed in a robot's manipulator in a known grasp.
 - (b) A base pose for the grasp is provided.
 - (c) The robot collects view from many different precisely known poses, together with models for self and background filtering.

Where k might be around 3.

4.5 Automatic Teaching

1. Place an object in front of the robot.
2. Robot repeats k times:
 - (a) Infer a stateless grasp on the object.
 - (b) **if** a pose model exists, infer the base pose for the grasp.
 else, start a new pose model.
 - (c) Perform semi-automatic teaching with the object.
 - (d) Replace the object.

Where k might be around 3.

5 Employing the System

Object Poses on Table or Floor

Single Target Tracking

6 Yet To Be Implemented

Background Filtering

Depth Models

Train feature weights to minimize the leave-one-out error rate on training data.

7 Attendance Statement

I have been paying attention to AI, Machine Learning, and Computer Vision since 2004. Regarding Computer Vision, I saw the progress of Neural Nets in the 90's swept under the rug by SIFT, HoG, and SVMs in the mid 2000's, only for neural nets to reclaim the throne in the 2010's. At one time it was said that AI was "vision hard" and that solving Vision would effectively solve AI. While that belief sweeps a bit under the rug, it is certainly true that in the past Computer Vision has been a strong bottleneck in the development of AI and Robotics. Recent advances in object detection on large data sets suggest that we are ready to move beyond attacking vision in an isolated setting and begin integrating it in a larger framework for planning in an interactive environment. I have training in state of the art computer vision techniques and have been tracking the literature for a few years now. Although I have attended CVPR twice, I have not attended any conferences in AI or Robotics. I recently started to focus on Real Time Vision systems, and attending AAAI will help me rapidly learn about the community and hit the ground running. One of my goals is to make a real time vision system with a capacity of 50 objects. It should be capable of giving accurate 3D pose estimates that enable the objects to be identified and manipulated by robots. It should be able to learn a new object in less than a minute, and should be simple enough that a non-expert can easily train the system. Yet it should remain versatile enough that state of the art detectors running on special hardware can be swapped in by experts. I have made good progress towards this goal so far. By attending AAAI, I will be able to better understand the motivations, needs, and desires of AI and Robotics concentrators with regards to Computer Vision, which will enable me to complete my dissertation in a way that aligns with these communities' values.

8 Curriculum Vitae

Education Florida State University, 2003-2006 UC Berkeley, 2006-2008 U Chicago, 2010-2011

Brown University, 2011-Present

Employment Havok, 2008-2009

Conference Papers CVPR NIPS

Conferences Attended CVPR colorado springs CVPR providence

Misc REU at OSU

9 Letter From Supervisor